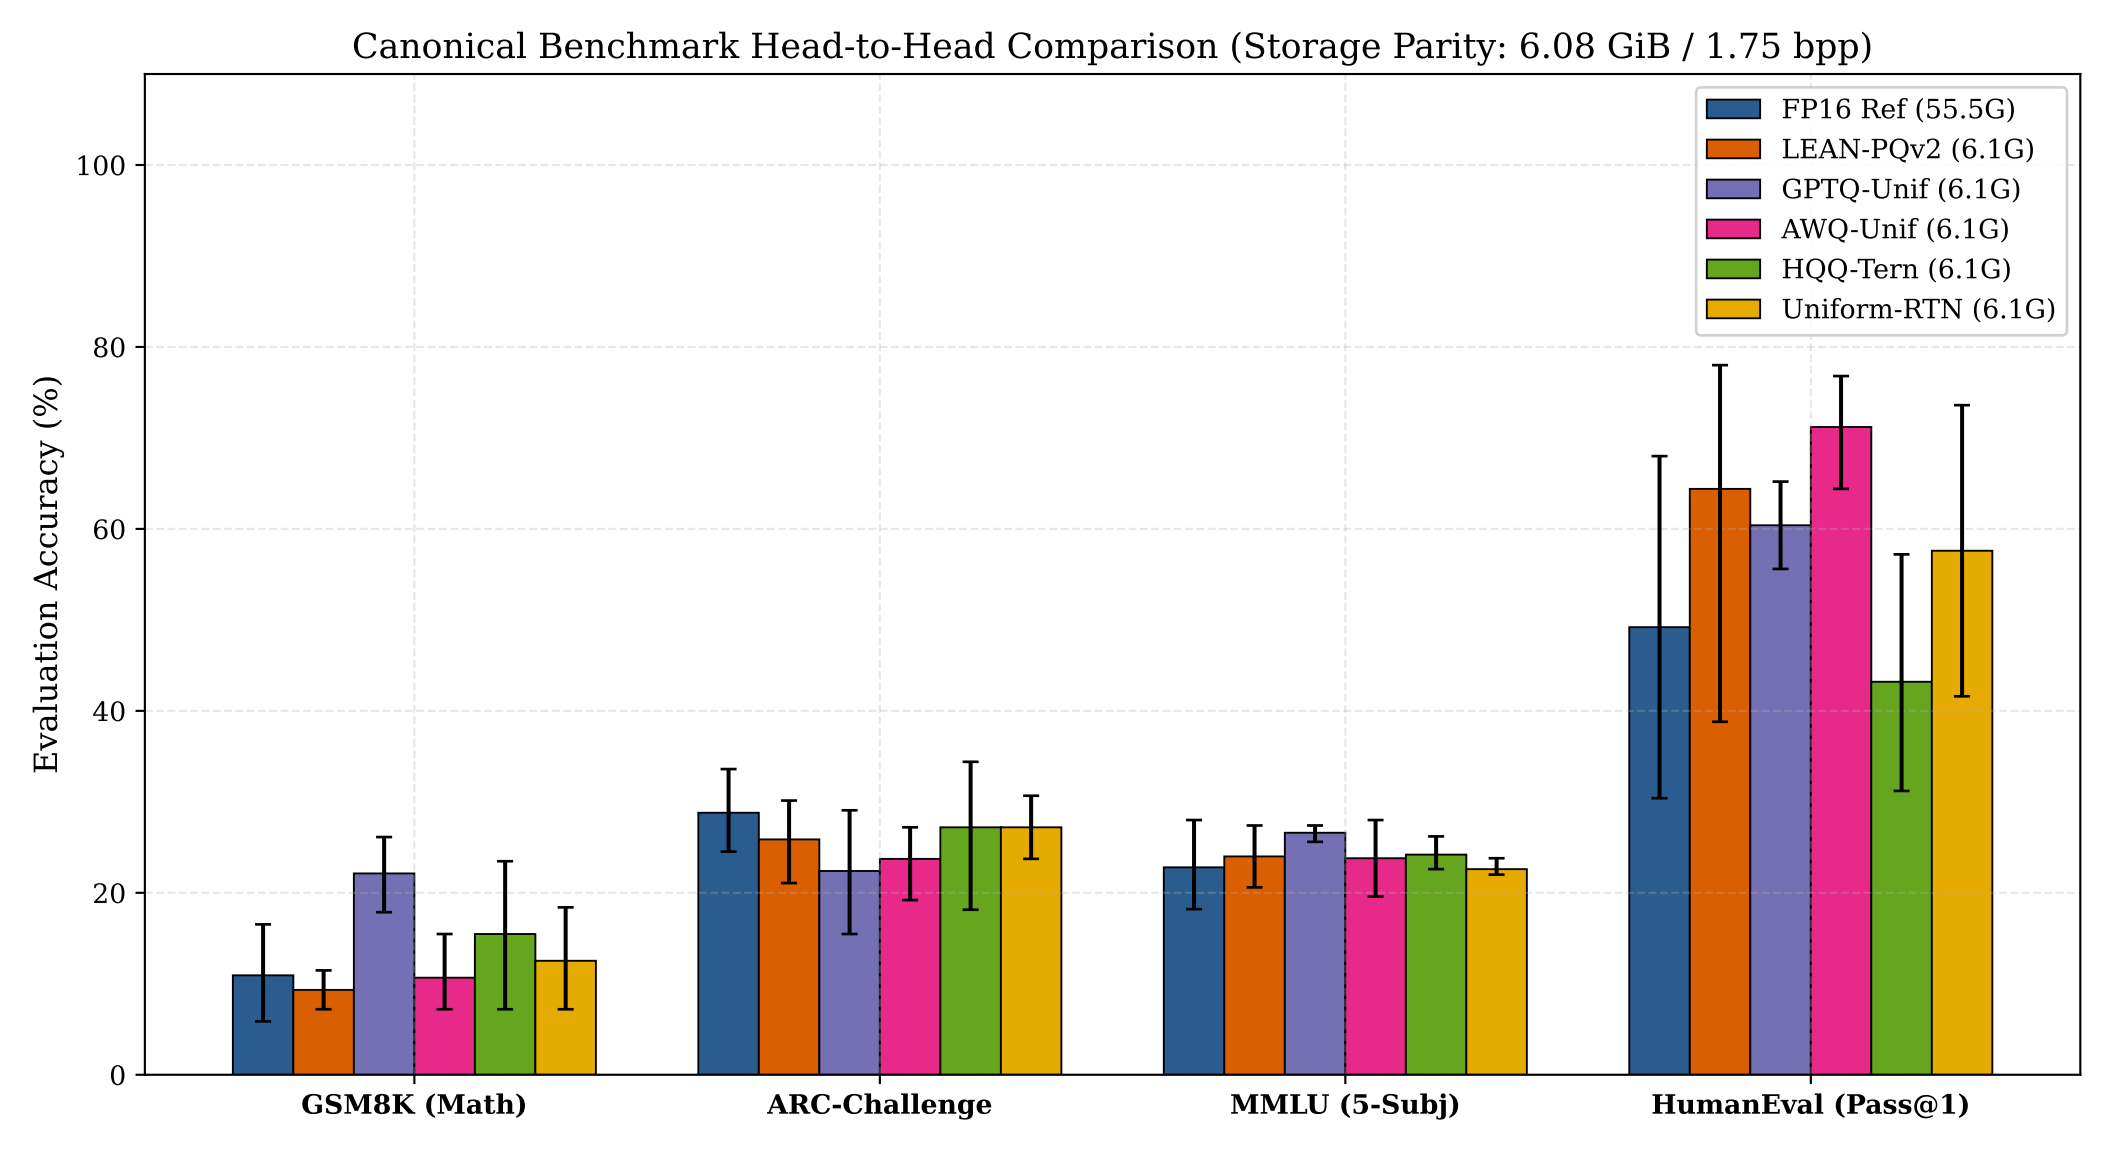
Canonical Benchmark Head-to-Head Comparison (Storage Parity: 6.08 GiB / 1.75 bpp)
Evaluation Accuracy (%)
0
20
40
60
80
100
FP16 Ref (55.5G)
LEAN-PQv2 (6.1G)
GPTQ-Unif (6.1G)
AWQ-Unif (6.1G)
HQQ-Tern (6.1G)
Uniform-RTN (6.1G)
GSM8K (Math)
ARC-Challenge
MMLU (5-Subj)
HumanEval (Pass@1)