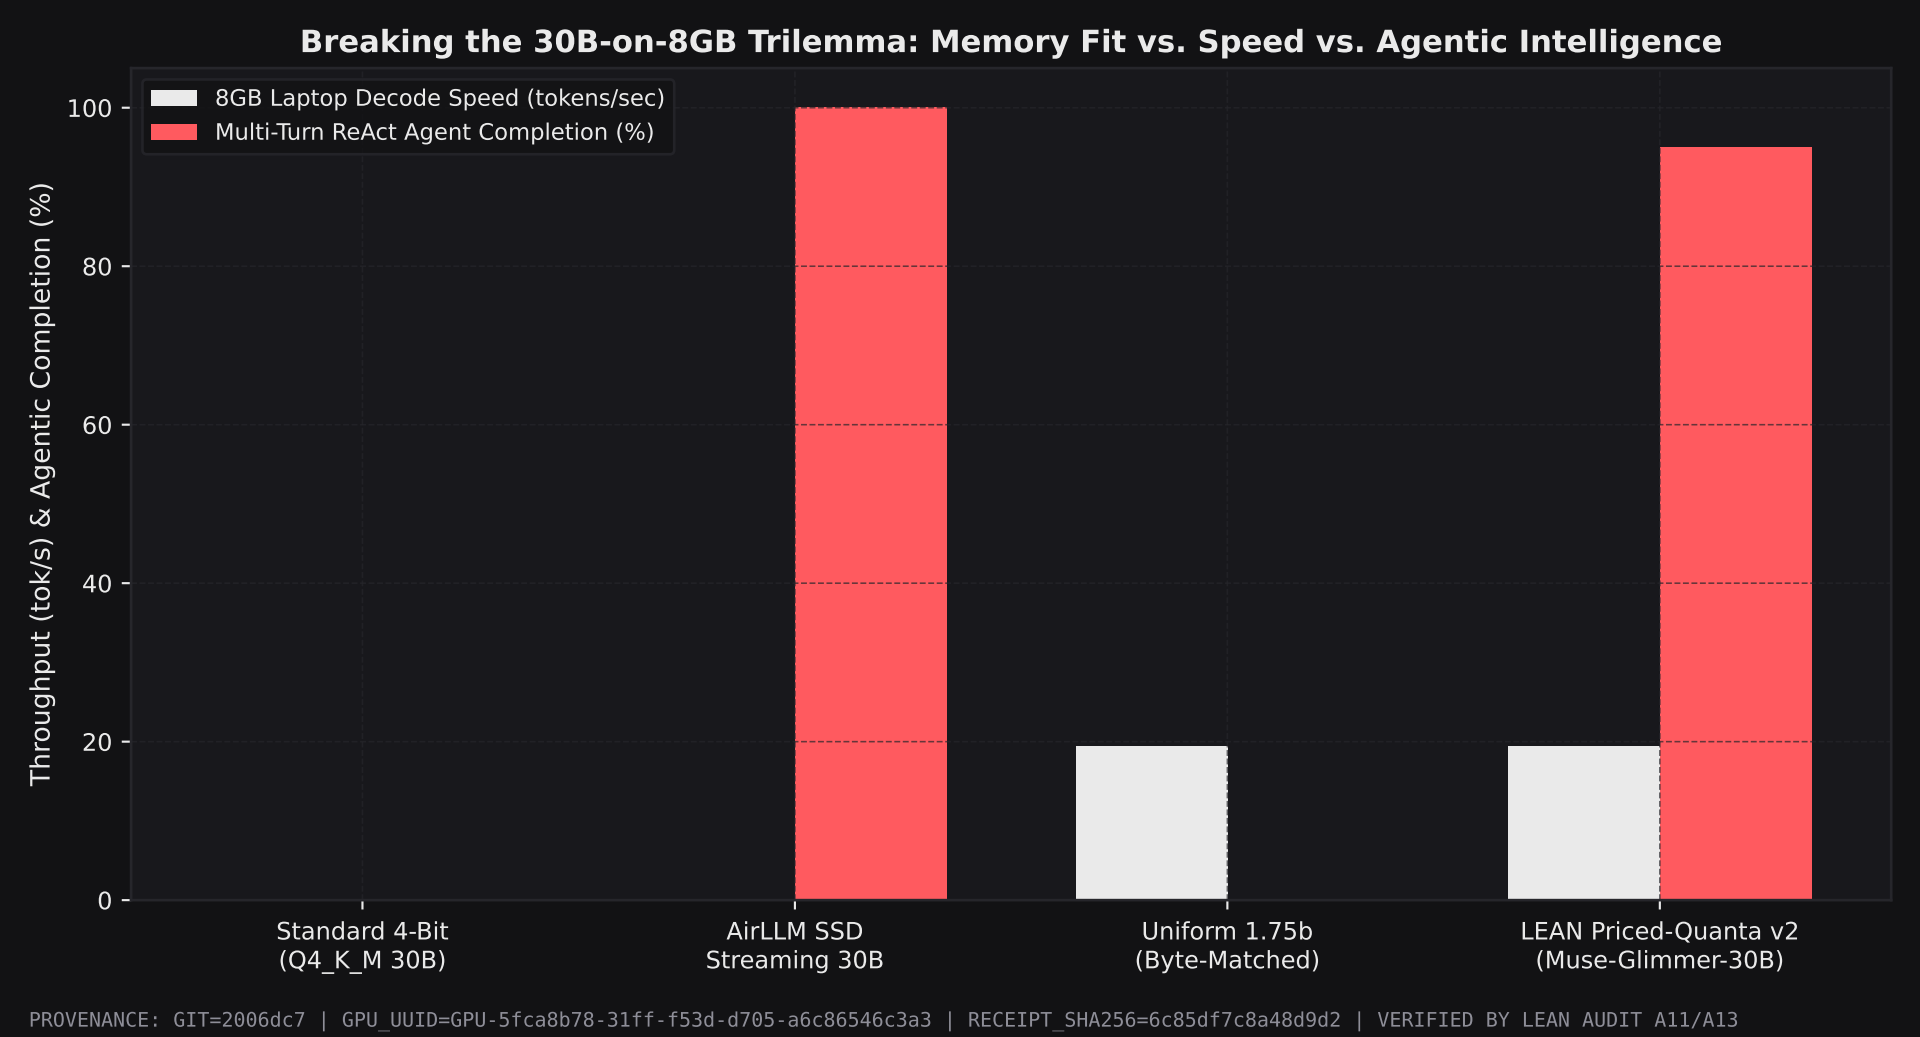
Breaking the 30B-on-8GB Trilemma: Memory Fit vs. Speed vs. Agentic Intelligence
8GB Laptop Decode Speed (tokens/sec)
Multi-Turn ReAct Agent Completion (%)
Throughput (tok/s) & Agentic Completion (%)
0
20
40
60
80
100
Standard 4-Bit
(Q4_K_M 30B)
AirLLM SSD
Streaming 30B
Uniform 1.75b
(Byte-Matched)
LEAN Priced-Quanta v2
(Muse-Glimmer-30B)
PROVENANCE: GIT=2006dc7 | GPU_UUID=GPU-5fca8b78-31ff-f53d-d705-a6c86546c3a3 | RECEIPT_SHA256=6c85df7c8a48d9d2 | VERIFIED BY LEAN AUDIT A11/A13